



University of New Haven

TAGLIATELA COLLEGE OF ENGINEERING

Electrical & Computer Engineering and Computer Science

## Electrical & Computer Engineering & Computer Science (ECECS)

# TECHNICAL REPORT



**Spring - 2024**

Cover Page	01
Contents	02
Summary	03
Title of Project	04
Highlights of the Project	04
Introduction	05
What is RFM?	05
Dataset	06
Data Pipeline	06
Methodology	07
CRISP – DM	07
Azure Data Factory	08
Databricks Integration	09
RFM Score Calculation	10
Code Implementation	11
Results	15
Conclusion	16

## Executive Summary

Using RFM (Recency, Frequency, Monetary) analysis for customer segmentation is an effective way to separate out various consumer categories according to how they transact. Businesses may process huge volumes of transactional data efficiently by using Spark for scalable and distributed RFM score calculation, and Azure Data Factory for data collection, preparation, and orchestration. Consumers can be divided into groups like High-Value, At-Risk, and Low-Value consumers after RFM ratings are determined. This enables firms to adjust their marketing plans and product offerings appropriately. In order to better facilitate the implementation of customized marketing campaigns aiming at enhancing consumer engagement and retention, Azure Data Factory may further coordinate the analysis of each segment's attributes and behavior. By using this strategy, companies may maximize client lifetime value and provide individualized experiences that improve overall marketing and sales effectiveness.

## Team Members

- **Siddhant Alhat Rajendra**
- **Veda Samohitha Chaganti**
- **Krishna Chaitanya Vutukuru**
- **Sathish Gandhi Parasuram Reddy**

# Title

## Title of Project

### Smart Segmentation

Harnessing Data Engineering for Customer Segmentation

## Highlights of Project

Using both Azure and Databricks to develop our data pipeline, the RFM (Recency, Frequency, Monetary) analysis project was essential to improving our understanding of customer behavior and preferences. With the help of Databricks' potent analytics platform and Azure's sturdy architecture, we were able to gather, prepare, and analyze transactional data with ease, enabling thorough insights. Using the scalable computing power of Databricks, we were able to determine important aspects of consumer behavior by calculating RFM scores. Future development prospects encompass investigating sophisticated segmentation strategies and incorporating supplementary data sources to achieve a more all-encompassing perspective of client conduct. In conclusion, the RFM analysis project—which is made possible by Azure and Databricks—represents a crucial step in the direction of maximizing marketing and customer relationship management tactics, which will promote company expansion and success.



**Submitted on: 22<sup>nd</sup> April 2024**

# Introduction

## Introduction

Customer segmentation helps with focused marketing and customer satisfaction by grouping customers according to common characteristics. This project collects, saves, and analyzes consumer data at scale using Apache Spark distributed data processing and Azure cloud architecture. Businesses are able to improve customer experiences, optimize resources, and refine marketing tactics through the use of Spark's clustering algorithms, which segment customers.

## What is RFM?

Three metrics are used by RFM, a customer segmentation technique, to assess the behavior of its customers: monetary, frequency, and recency. It considers the amount, frequency, and recentness of a customer's purchases. Through the integration of these variables, companies can discern advantageous consumer groups, facilitating focused advertising, customized correspondence, and efficient resource distribution to enhance customer involvement and allegiance.

## Importance of RFM Analysis

- **Widely adopted:** Major companies across industries use RFM analysis for marketing campaigns, customer relationship management, and business decision-making.
- **Benefits for businesses:** RFM analysis helps businesses identify high-value customers, tailor marketing strategies, improve customer retention, and optimize resource allocation.

# Data Pipeline

## Dataset


This is a transnational data set which contains all the transactions occurring between 01/12/2010 and 09/12/2011 for a UK-based and registered non-store online retail. The company sells unique all-occasion gifts. Many customers of the company are wholesalers.

## Data Pipeline


We used Azure to create our data pipeline since it provided a stable platform that met our demands. Utilizing Azure's services improved our productivity and gave us dependable resources for efficient data management. We could leverage Azure's scalability and integration features to make sure our data pipeline was effective and flexible enough to change as our needs did.



Data ingestion through  
GA tags (Extracting the  
data)



Data preprocessing in  
Azure



Calculating RFM score  
in Databricks and  
implementing spark

# Data Pipeline

This screenshot shows the Microsoft Azure Data Factory portal interface. The main canvas displays a data pipeline named 'pipeline1' with a single activity 'Data flow1' connected to a 'Notebook' activity 'Notebook1'. The left sidebar shows the 'Activities' list with categories like 'Move and transform', 'Azure Data Explorer', 'Databricks', 'Data Lake Analytics', and 'General'. The right sidebar shows the 'Properties' panel for the selected activity, with fields for 'Name' (pipeline1) and 'Description'. The bottom status bar shows the system clock as 03:31 on 22-04-2024.

This screenshot shows a more complex data pipeline in the Microsoft Azure Data Factory portal. The main canvas displays a data pipeline named 'pipeline1' with a sequence of activities: 'source1', 'CountryUK', 'Pricegreatertha...', 'Quantitygreater...', 'filter1', and 'sink1'. There are also branching activities like 'NotUK' and 'Leastthanzero'. The left sidebar shows the 'Factory Resources' list, including 'Pipelines', 'Change Data Capture (preview)', 'Datasets', 'Data flows', and 'Power Query'. The right sidebar shows the 'Properties' panel for the selected activity, with a 'Related (1)' section listing 'pipeline1'. The bottom status bar shows the system clock as 5:44 AM on 4/22/2024.



# Methodology

## CRISP - DM

We used the CRISP-DM (Cross-Industry Standard Process for Data Mining) methodology for our study since it offered an organized way to direct our analysis. We used CRISP-DM in the following ways for our RFM study project.

**Business Understanding:** During the first stage, we concentrated on fully comprehending the dataset as well as the goals of our study. we carefully examined the dataset to understand its composition, organization, and insights. Our goal was to provide significant results that advance the field of marketing analytics by comprehending the context of the information and its applicability to the analysis of consumer behavior.

**Data Understanding:** We then went into detail to comprehend the dataset. To do this, we had to go through the dataset and find its features, their distributions, and any obvious patterns. Additionally, we evaluated the quality of the data, looking for anomalies, missing numbers, and consistency issues. Our choices for data pretreatment and modeling techniques were informed by this in-depth knowledge of the dataset.

**Data Preparation:** After gaining a thorough grasp of the dataset, we moved on to the data's preparation for analysis. This included fixing mistakes or inconsistencies in the data, changing variables as necessary, and, where relevant, merging data from several sources. Furthermore, to ensure reliable and significant results, we formatted the data in a way that made it ideal for RFM analysis.

**Modeling:** To begin the modeling stage, we determined each customer's RFM scores. We used methods and algorithms to calculate frequency, recency, and



# Azure

monetary value metrics based on transactional data by utilizing tools like Spark. Effective client segmentation and the discovery of useful insights depended on this stage.

## Azure Data Factory

- **Role in basic data processing:** Azure Data Factory was used to construct a pipeline for basic data processing and transformations.
- **Construction of pipeline:** The pipeline included steps for data ingestion, cleansing, and transformation using Azure Data Flow.
- **Key steps in data transformation:** Data was filtered, aggregated, and prepared for further analysis in Databricks.

The screenshot displays the Azure Data Factory (ADF) portal interface. The top navigation bar shows the 'Data Factory' tab selected. The left sidebar contains 'Factory Resources' with a tree view showing 'Pipelines' > 'pipeline1'. The main canvas shows the 'Activities' section for 'pipeline1', which includes a 'Data flow' activity (Data flow1) and a 'Notebook' activity (Notebook1). The 'Data flow' activity is currently selected, and its 'Output' tab is active, showing a table with one row of data. The 'Properties' panel on the right shows the 'General' tab for the selected activity, with fields for 'Name' (pipeline1) and 'Description'. The bottom status bar indicates the pipeline is 'In progress' with a run ID of 'ab64bdd1-ce0d-471e-968c-97e6e73bc2bc' and a run start time of '4/22/2024, 2:05:17'.

# Databricks

## Databricks Integration

- **Utilization of Spark code:** Spark code was implemented in Databricks for advanced data processing tasks.
- **Filtering null values:** Null values were filtered out from the dataset to ensure data quality.
- **Calculating RFM scores:** RFM scores were calculated using Spark code to evaluate the recency, frequency, and monetary value of customer transactions.

The screenshot displays the Databricks web interface within a Microsoft Azure browser window. The notebook, titled 'Untitled Notebook 2024-04-22 00:48:01', is running on the 'Alhat, Rajendra's Cluster'. The interface includes a left sidebar with navigation options like 'New', 'Workspace', 'Catalog', 'Workflows', 'Compute', 'SQL', 'Alerts', 'Query History', 'SQL Warehouses', 'Data Engineering', 'Job Runs', 'Data Ingestion', 'Delta Live Tables', 'Machine Learning', 'Playground', and 'Experiments'. The main area shows a notebook with three code cells:

- Cell 1: Spark configuration to retrieve cluster information.
 

```
spark.conf.get("spark.databricks.clusterUsageTags.clusterName")
spark.conf.get("spark.databricks.clusterUsageTags.clusterId")
```

 The output shows the cluster name '0422-041613-7jba4tw0'.
- Cell 2: Setting up connections and credentials.
 

```
#Connections credentials
ContainerName = "output"
azure_blobstorage_name = "teamsde"
mountpointname = "/mnt/rfm"
secret_key = "15w1sV649yQnqL40e1sOGL1PG1Qq8m3HitaqjydN2QtUk8bYUk13rEXqYnrTuCZSD78Wip5HA3E+ASTeYpsgA="
```
- Cell 3: Mounting the storage.
 

```
#mount
#dbutils.fs.mount(source = f"wasbs://{output@teamsde.blob.core.windows.net}", mount_point = mountpointname ,
#extra_configs = {"fs.azure.account.key":"azure_blobstorage_name+.blob.core.windows.net":secret_key})
```

Below the code cells, there is a section for 'Skipped' commands. An assistant chat window on the left provides guidance on fixing an error related to file modes, suggesting the use of 'overwrite' when saving a DataFrame as a CSV file.

# RFM Score

## RFM Score Calculation

- **Explanation of RFM score components:** Recency, Frequency, and Monetary (RFM) scores are calculated based on the customer's recent purchase, frequency of purchases, and monetary value of transactions.
- **How RFM scores are calculated:** Each customer is assigned an RFM score based on their transaction history, with higher scores indicating higher value or engagement.

The screenshot shows a Databricks notebook titled "Untitled Notebook 2024-04-22 00:48:01". The notebook is running Python code to calculate RFM scores. The code uses the `r_f_m` function to calculate the scores and then displays the results using `r_f_m_seg.show(5)`.

The output shows a table with 7 columns: `CustomerID`, `Recency`, `Frequency`, `Monetary_value`, `r_score`, `f_score`, and `m_score`. The table displays the top 5 rows of results.

CustomerID	Recency	Frequency	Monetary_value	r_score	f_score	m_score
15727	16	7	5178.96	2	3	1
17389	0	34	31833.68	1	2	1
16503	106	4	1431.93	NULL	3	1
14570	280	2	218.06	NULL	4	2
17420	50	3	598.83	3	4	2

The notebook also includes an Assistant chat window on the left, which provides instructions on how to fix an error related to the `write.mode` parameter.

# Code Implementation

This section describes the Databricks code implementation used to conduct RFM (Recency, Frequency, Monetary) analysis. We have streamlined the data ingestion, preparation, and analysis process to extract useful insights from our transactional data by utilizing the capabilities of Databricks. The procedures for determining RFM scores. Our goal with using Databricks is to maximize the effectiveness and scalability of our RFM analysis, which will improve our comprehension of customer segmentation and allow for more focused marketing approaches.

```
#Connections credentials
ContainerName = "output"
azure_blobstorage_name = "teamdsde"
mountpointname = "/mnt/rfm"
secret_key
="15wIsV649yQmql40e1sOGLlPGlQq8mJHitaqjydN2q0tUkBbYUXiJrEXqYnrTuCZSDJ8Wip5H
A3E+AStYEpsgA=="
```

```
#mount
dbutils.fs.mount(source =
f"wasbs://output@teamdsde.blob.core.windows.net",mount_point =
mountpointname ,extra_configs =
{"fs.azure.account.key."+azure_blobstorage_name+".blob.core.windows.net":se
cret_key})
```

```
DF = spark.read.format('csv').options(
    header='true', inferSchema='true').load("/mnt/rfm/*.csv")
```

```
rowcount=DF.count()
colcount=len(DF.columns)
print("Dataframe rows:",rowcount)
print("Dataframe cols:",colcount)
```

```
#Null count
from pyspark.sql.functions import col,isnan, when, count
```

# Code Implementation

```
DF.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in
DF.columns]
).show()
```

```
DF1=DF.na.drop(subset=["CustomerID"])
DF1.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in
DF1.columns]
).show()
```

```
DF1.printSchema()
DF1.show(10)
```

```
#Imports
# #def square(x):
#     return x**2

import pandas as pd
pandasDF = DF1.toPandas()
pandasDF['InvoiceDate'] = pd.to_datetime(pandasDF['InvoiceDate'])
pandasDF["Date"] = pd.DatetimeIndex(pandasDF["InvoiceDate"]).date
pandasDF["Month"] = pandasDF["InvoiceDate"].dt.month
pandasDF["Year"] = pandasDF["InvoiceDate"].dt.year
DF2 = spark.createDataFrame(pandasDF)
DF2.show()
```

```
from pyspark.sql.functions import *
from pyspark.sql.types import *
import pyspark.sql.functions as F
from pyspark.sql.types import *
DF3=DF2.withColumn("Totalsales",col("UnitPrice")*col("Quantity"))
Totalsalesbymonth=DF3.groupBy('Description').sum('Totalsales')
Totalsalesbymonth.show()
```

```
date_max = DF3.select(F.max(DF3.InvoiceDate).alias('max_date')).toPandas()
DF4 = DF3.withColumn('Duration', F.datediff(F.lit(date_max.iloc[0][0]),
'InvoiceDate')) #'2011-12-04 13:15:00'
DF4.show(5)
```

# Code Implementation

```
recency = DF4.groupby('CustomerID').agg(F.min('Duration').alias('Recency'))
recency.sort(col('Recency')).show()
```

```
frequency = DF4.groupby('CustomerID', 'InvoiceDate').count()\
                .groupby('CustomerID')\
                .agg(F.count("*").alias("Frequency"))
frequency.orderBy(frequency.Frequency.desc()).show(5)
```

```
monetary = DF4.groupby('CustomerID').agg(F.round(F.sum('Totalsales'),
2).alias('Monetary_value'))
monetary.orderBy(F.col('Monetary_value').desc()).show(5)
```

```
rfm = recency.join(frequency, on = 'CustomerID', how = 'inner')\
            .join(monetary, on = 'CustomerID', how = 'inner')
rfm.show(5)
```

```
rfm.where(rfm.Recency.contains(0)).show()
```

```
def r_score(r_value):
    """recency with lower value ranks the least, i.e the
    more recent a user transaction, the lower the rank"""

    if r_value <= 14:    # 2 weeks
        return 1
    elif r_value <= 31:  # 1 month
        return 2
    elif r_value <= 93:  # 3 month
        return 3
    else:
        return

def f_score(f_value):
    """the more a customer transact with the business,
    the more their rank is closer to 1, the fewer the number of
    times they have made transaction, the higher the rank"""

    if f_value <= 3:
        return 4
```

# Code Implementation

```
elif f_value <= 18:
    return 3
elif f_value <= 36:
    return 2
else:
    return 1
```

```
def m_score(m_value):
    """the more the value contributed to the business,
    the higher their rank to 1"""

    if m_value < 10:
        return 4
    elif m_value <= 100:
        return 3
    elif m_value <= 1000:
        return 2
    else:
        return 1
```

```
r_udf = F.udf(lambda r_value: r_score(r_value), StringType())
f_udf = F.udf(lambda f_value: f_score(f_value), StringType())
m_udf = F.udf(lambda m_value: m_score(m_value), StringType())
```

```
rfm_seg = rfm.withColumn('r_score', r_udf(F.col('recency')))
rfm_seg = rfm_seg.withColumn('f_score', f_udf(F.col('frequency')))
rfm_seg = rfm_seg.withColumn('m_score', m_udf(F.col('monetary_value')))
rfm_seg.show(5)
```

```
rowcount=rfm_seg.count()
colcount=len(rfm_seg.columns)
print("Dataframe rows:",rowcount)
print("Dataframe cols:",colcount)
```

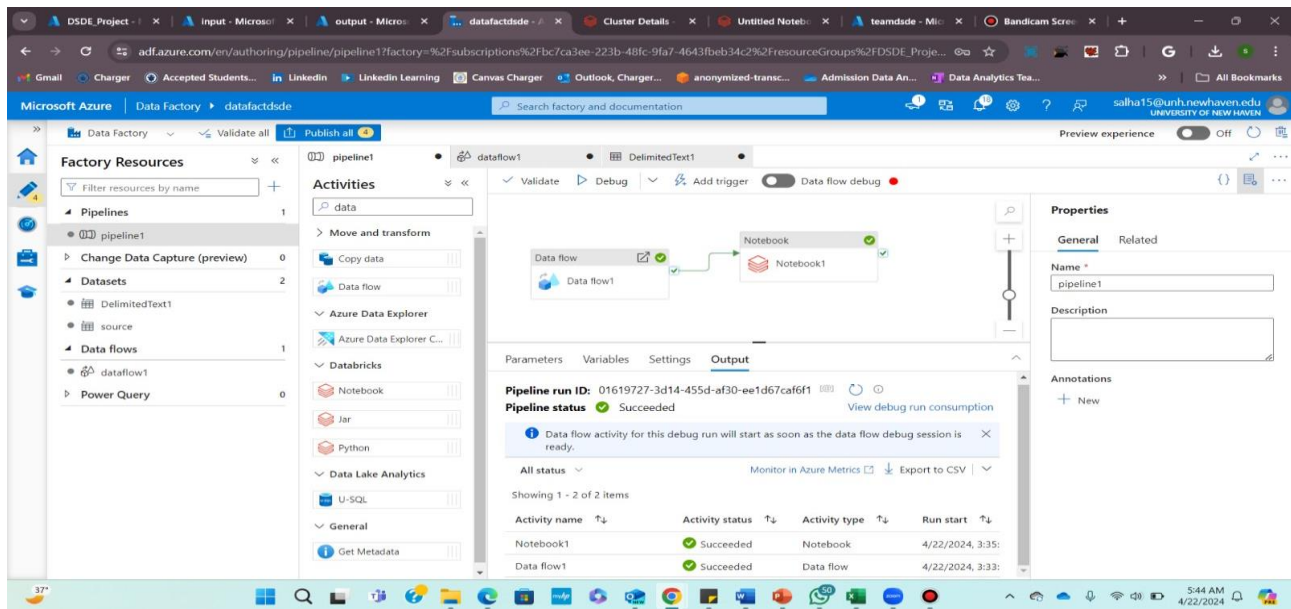
```
rfm_seg.write.mode("overwrite").format("csv").save("/mnt/RFMdata/RFMresults/RFMscores.csv")
```



# Results

## Results

The use of Azure to create our data pipeline produced important insights, most notably RFM scores. Through improved data analysis, we produced RFM (Recency, Frequency, Monetary) scores that offer a more comprehensive view of consumer behavior and preferences. One important outcome of our Azure-based data pipeline solution is the ability to compute RFM scores, which shows the real-world advantages of using advanced data analytics methods in our project.



```
rfm_seg: pyspark.sql.dataframe.DataFrame = [CustomerID: integer, Recency: integer ...
```

CustomerID	Recency	Frequency	Monetary_value	r_score	f_score	m_score
15727	16	7	5178.96	2	3	1
17389	0	34	31833.68	1	2	1
16503	106	4	1431.93	NULL	3	1
14570	280	2	218.06	NULL	4	2
17420	50	3	598.83	3	4	2

## Customer Segmentation and Insights

- **Application of RFM scores:** RFM scores are used to segment customers into diverse groups based on their buying behavior.
- **Interpretation of RFM scores:** Businesses can gain insights into customer behavior and preferences, enabling targeted marketing campaigns and personalized customer experiences.
- **Business insights and strategies:** RFM analysis helps businesses identify high-value customers, re-engage dormant customers, and optimize marketing strategies for better ROI.

# Conclusion

## Conclusion

Our project's goals, which are to improve customer experiences and optimize marketing strategies, depend heavily on consumer segmentation utilizing Spark and Azure. We aim to explore the complexities of RFM analysis, a popular technique for classifying clients based on recency, frequency, and monetary value, by utilizing the power of these innovative technologies. Our project's main goal is to effectively handle large amounts of transactional data by utilizing Azure's cloud infrastructure, which has strong data processing and storage capabilities. In the meantime, real-time analysis and faster calculation times are made possible by Spark's distributed computing framework, which allows for parallel processing across clusters. Our goal is to discover high-value client groupings and adjust marketing campaigns accordingly by using the synergy between Spark and Azure. ultimately generating prospects for corporate growth and profitability.

Through our project, we hope to demonstrate how Azure and Spark can revolutionize the fields of marketing analytics and customer segmentation. Through resource optimization and the utilization of innovative data processing skills, our goal is to show how companies can extract meaningful insights from massive datasets, which can enhance customer happiness and better target marketing campaigns. We aim to enable organizations to uncover new revenue streams, improve competitiveness, and prosper in today's changing marketplace through our thorough analysis and strategic execution.